

OS 2025 Lab 3

Multithreading Program & Linux Kernel Module

Due Date: **2025/12/26 17:00 (before lab3 course finish)**



TA: 陳煒勛、徐健翔、鄭煦霖


Outline

- A. Objectives
- B. Background & Requirements
 - 1. Lock
 - 2. Multithread
 - 3. Kernel Module
- C. Grading

Objectives

- **Process Synchronization**
 - Learn how to protect critical section
- **Multithreading**
 - Take advantage of multi-core systems
 - Beware potential synchronization problem
- **Linux Kernel Module with Proc File System**
 - Communicate with kernel using Proc file

Race Condition



Race Condition

Scenario 1

thread1

Load a //a=0
Add a 1 //a=1
Store a //a=1

thread2

Load a //a=0
Add a 1 //a=1
Store a //a=1

Output: a=1

Scenario 2

thread1

Load a //a=0
Add a 1 //a=1
Store a //a=1


thread2

Load a //a=1
Add a 1 //a=2
Store a //a=2

Output: a=2

Time

Race Condition



How to protect critical section ?



```
do  
{  
    entry section  
    critical section //update shared data here...  
    exit section  
    remainder section  
} while (TRUE)
```

Pthread Spin lock

- pthread_spin_lock
 - The `pthread_spin_lock()` function locks the spin lock referred to by lock. If the spin lock is currently unlocked, the calling thread acquires the lock immediately.
- Pthread_spin_unlock
 - The `pthread_spin_unlock()` function unlocks the spin lock referred to lock. If any threads are spinning on the lock, one of those threads will then acquire the lock.

Requirement – 1.1

Assignment 1.1

- You will get ***1_1.c, 1_ans.txt, judge.out, and Makefile*** in Assignment 1.1
- Fill ***/*YOUR CODE HERE*/*** in ***1_1.c*** .
- Test your code with Makefile.

Assignment 1.1

Two threads which increment a by 1 are given.

Protect the critical section with spin lock.

Expectation: a=20000

Use Makefile to validate your code!

```
iloveos@iloveos-VirtualBox:~/os_hw/answer/1/1_1$ make  
Success
```

Complete your code insides this area!

```
/*YOUR CODE HERE*/  
critical section  
*****
```

Spin Lock

- Atomic operation XCHG(exchange)
 - XCHG instructions swaps the contents of two operands.
 - When a memory operand is used with the XCHG instruction, the processor's LOCK signal is automatically asserted.

Atomic Instruction

- **xchg**
 - Exchange values of register or memory
 - **Atomic** operation

XCHG Instruction




XCHG exchanges the values of two operands. At least one operand must be a register. No immediate operands are permitted.

```
.data
var1 WORD 1000h
var2 WORD 2000h
.code
xchg ax,bx      ; exchange 16-bit regs
xchg ah,al      ; exchange 8-bit regs
xchg var1,bx    ; exchange mem, reg
xchg eax,ebx    ; exchange 32-bit regs

xchg var1,var2  ; error 2 memory operands
```


Spin Lock

- Atomic operation
 - xchg



Spin Lock

- Atomic operation
 - xchg



Requirement – 1.2

Assignment 1.2

- You will get ***1_2.c, 1_ans.txt, judge.out, and Makefile*** in Assignment 1.2
- Fill ***/*YOUR CODE HERE*/*** in ***1_2.c***.
- Test your code with Makefile.

Requirement – 1.2

Assignment 1.2

Two threads which increment a by 1 are given.

Protect the critical section with spin lock.

Using **xchg** instruction to complete the spin lock by yourself.

Expectation: a=20000


```
#define LOCK 0  
#define UNLOCK 1
```

Use Makefile to validate your code!

```
iloveos@iloveos-VirtualBox:~/os_hw/answer/1/1_1$ make  
Success
```


Multithreading Program

Single thread



Assignment 2.1

Two threads



Assignment 2.2

Your work!

Beware race condition problem!

Assignment 2

- You will get ***2_1.c, 2_2.c, m1.txt, m2.txt and Makefile*** in Assignment 2.1 and Assignment 2.2.
- Fill ***/*YOUR CODE HERE*/*** in ***2_1.c and 2_2.c***.
- Test your code with Makefile.

Requirement – 2.1 & 2.2

Assignment 2

2.1

Complete the matrix multiplication with single thread.

2.2

Complete the matrix multiplication with two threads.

Use Makefile to validate your code!

2.1


```
iloveos@iloveos-VirtualBox:~/os_hw/answer/2$ make judge1
Success
```

2.2

```
iloveos@iloveos-VirtualBox:~/os_hw/answer/2$ make judge2
Success
```

Linux Kernel Module

- The linux kernel has the ability to **load** and **unload** arbitrary sections of kernal code on demand. These loadable kernel modules run in priviledged kernal mode. In theory, there is no restriction on what a kernael module is allowed to do.
- Linux kernel can load and unload modules dynamically at run time. It saves recompiling, relinking and reloading time.



Follow the guide. Try to write a HelloWorld first by yourself!

Linux Kernel Module(In This Assignment)

- Program
 - Create 2 threads like Assignment 2.1 and 2.2.
- Linux Kernel Module
 - Create **/proc** file.
 - Complete the **Read** operation and **Write** operation.
 - Show the process information through **Read** and **Write** operations.

Linux Process File System

- Pseudo file system -> Data is not stored persistently.
- Provide a way for user programs to access process information as plain text files.
- For example, in the past, traditional UNIX **ps** command has been implemented as a privileged process that reads the process state directly from kernel's virtual memory. Under Linux, this command is implemented as an entirely unprivileged program that simply parses the information from /proc.

Linux Process File System

- `proc_create(name, mode, parent, proc_ops)`
 - Create a proc file under proc file system.

```
iloveos@iloveos-VirtualBox:/proc$ ls | grep My  
Mythread_info
```

- Struct `proc_ops`
 - We use 2 operations(function pointers) in this struct.
 1. `proc_read`
 2. `proc_write`

Read & Write


- When read(write) proc file, function proc_read(proc_write) will be executed.
- Therefore, our objectives are to complete the desired read and write operation.

Access Kernel

- `copy_to_user`: copies n bytes from the kernel-space, from the address referenced by `from` in user-space.
- `copy_from_user`: copies n bytes from user-space from the address referenced by `from` in kernel-space.


Read

- `ssize_t read(file, ubuf, size, offset)`
- The module is responsible for **advancing the offset** according to how much it reads and **returning the read size**.



Write


- `ssize_t write(file, ubuf, size, offset)`
- The module is responsible for **advancing the offset** according to how much it reads and **returning the write size**.



Assignment 3

- You will get ***3_1.c, My_Kernel.c, m1.txt, m2.txt and Makefile*** in Assignment 3.1 and ***3_2.c, My_Kernel.c, m1.txt, m2.txt and Makefile*** in Assignment 3.2.
- Assignment 3.1:
Fill ***/*YOUR CODE HERE*/*** in ***My_Kernel.c***
- Assignment 3.2:
Fill ***/*YOUR CODE HERE*/*** in ***3_2.c My_Kernel.c***
- Test your code with Makefile.

Overall Flow(Assignment 3.1)



Assignment 3.1

- Through proc_read, Show process ID, thread ID, priority and state.


```
iloveos@iloveos-VirtualBox:~/os_hw/answer/3_1$ make Prog
PID: 2755, TID: 2756, Priority: 120, State: 0
PID: 2755, TID: 2757, Priority: 120, State: 0
```

Before make Prog:

current->pid	for_each_thread(current, thread)	thread->pid	thread_prio	thread->__state
--------------	----------------------------------	-------------	-------------	-----------------

1. make all
Build the modules.
2. make load/unload
Load/unload the module to kernel.

Overall Flow(Assignment 3.2)



Assignment 3.2

1. Through proc_read and proc_write, show string, process ID, thread ID and time(ms) in 1 thread.

```
iloveos@iloveos-VirtualBox:~/os_hw/answer/3_2$ make Prog_1thread
Thread 1 says hello!
PID: 47029, TID: 47030, time: 13600
current->tgid          current->utime/100/1000
                           current->pid
```

2. Through proc_read and proc_write, show string, process ID, thread ID and time(ms) in 2 threads.

```
iloveos@iloveos-VirtualBox:~/os_hw/answer/3_2$ make Prog_2thread
Thread 1 says hello!
PID: 47069, TID: 47070, time: 6360
Thread 2 says hello!
PID: 47069, TID: 47071, time: 6400
```

Before make Prog:

1. make all
Build the modules.
2. make load/unload
Load/unload the module to kernel.

Grading

- 1 Lock
 - 1.1 (1 points) pthread spin lock
 - 1.2 (1 points) write a spin lock
- 2 Multithreading(pthread)
 - 2.1 (0.5 points) matrix multiplication(1 thread)
 - 2.2 (1 points) matrix multiplication(2 threads)
- 3 Kernel module
 - 3.1 (1.5 points) proc file read
 - 3.2 (2 points) proc file read and write
- Answer 3 questions
 - Each question (1 points)

Reference

- Operating System Concepts 8th Edition
 - Chapter 4 Multithreaded Programming
 - Chapter 6 Synchronization
 - Chapter 21 The Linux System
 - 21.3 Kernel modules
 - 21.7.4 The Linux Process File System
- Intel® 64 and IA-32 Architectures Software Developer Manuals
- <https://sysprog21.github.io/lkmpg/>
- <https://en.wikipedia.org/wiki/Spinlock>
- https://linux-kernel-labs.github.io/refs/heads/master/labs/device_drivers.html