

# OS File System

OS 2025 LAB 4

DUE DATE: 2025/01/02 17:00  
(before lab 4 course finishes)

TA: NN6131037 陳彥廷、P76134846 王信智、P76144613 周安

# OUTLINE


1. Introduction to Virtual File System (VFS)

2. Layout of osfs

3. Access Path in File System


4. Requirement & Grading

# INTRODUCTION TO VFS



The VFS allows multiple file systems to coexist by separating generic functionality from specific implementations. The file system interface interacts with VFS, which selects and invokes the appropriate file system without needing to know their internal details.

# INTRODUCTION TO VFS



# STRUCTURE IN VFS

- Superblock
- Index node (inode)
- Directory entry (dentry)

# STRUCTURE IN VFS – SUPERBLOCK

- The superblock **stores the metadata of the filesystem**, describing its overall structure and state.
  - File system type (e.g., ext4, NTFS, etc.)
  - Total and available space
  - Data block size
  - Total number of files and directories

# STRUCTURE IN VFS – INODE

- An inode stores the attributes for a file or directory but does not store its name or content.
- Every file or directory has a corresponding inode that stores:
  - File size
  - Permissions (read/write/execute)
  - File type (regular file, directory, etc.) (i\_mode)
  - Pointers to the file's data blocks (location of actual file content)

# STRUCTURE IN VFS - DENTRY

- A dentry maps file or directory names to their corresponding inodes; it can be seen as an entry in a mapping table.
- Refer to Ch.14 p.16
- [Pathname lookup linux kernel document](#)

FILE NAME	INODE NUMBER
NAME1	13579
NAME2	13688

# VFS STRUCTURES IN OSFS

- Superblock
  - Index node (inode)
  - Directory entry (dentry)
- 
- osfs\_sb\_info
  - osfs\_inode
  - osfs\_dir\_entry

# OUTLINE


1. Introduction to Virtual File System (VFS)

2. Layout of osfs

3. Access Path in File System

4. Requirement & Grading

# Layout of OSFS




- Memory-virtualized disk

# AYOUT OF OSFS

```
struct osfs_sb_info {  
    uint32_t magic;          // Magic number to identify the filesystem  
    uint32_t block_size;     // Size of each data block  
    uint32_t inode_count;    // Total number of inodes  
    uint32_t block_count;    // Total number of data blocks  
    uint32_t nr_free_inodes; // Number of free inodes  
    uint32_t nr_free_blocks; // Number of free data blocks  
    unsigned long *inode_bitmap; // Pointer to the inode bitmap  
    unsigned long *block_bitmap; // Pointer to the data block bitmap  
    void *inode_table;        // Pointer to the inode table  
    void *data_blocks;        // Pointer to the data blocks area  
};
```


Superblock    Inode bitmap    Data bitmap    Inode table    Data blocks



# AYOUT OF OSFS

```
struct osfs_inode {  
    uint32_t i_ino;                      // Inode number  
    uint32_t i_size;                     // File size in bytes  
    uint32_t i_blocks;                   // Number of blocks occupied by the file  
    uint16_t i_mode;                    // File mode (permissions and type)  
    uint16_t i_links_count;             // Number of hard links  
    uint32_t i_uid;                     // User ID of owner  
    uint32_t i_gid;                     // Group ID of owner  
    struct timespec64 __i_atime;        // Last access time  
    struct timespec64 __i_mtime;        // Last modification time  
    struct timespec64 __i_ctime;        // Creation time  
    uint32_t i_block;                   // Simplified handling, single data block pointer  
};
```

Superblock    Inode bitmap    Data bitmap    Inode table    Data blocks



# OUTLINE

1. Introduction to Virtual File System (VFS)

2. Layout of osfs

3. Access Path in File System

4. Requirement & Grading


# ACCESS PATH IN FILE SYSTEM

- 
- We use `read()` as an example
  - Assume that we want to read data in "`/os/lab4`"

# ACCESS PATH IN FILE SYSTEM


## – *READ()*

1. Read the **root** inode and locate its data block pointers




# ACCESS PATH IN FILE SYSTEM – READ()

2. Read the **root** data block and locate the **os** entry



# ACCESS PATH IN FILE SYSTEM – READ()

3. Read the **os** inode and locate its data block pointers



Superblock

Inode bitmap

Data bitmap


Inode table

Data blocks

# ACCESS PATH IN FILE SYSTEM


## – READ()

4. Read the **os** data block and locate the **lab4** entry




# ACCESS PATH IN FILE SYSTEM – READ()

5. Read the **lab4** inode into memory




# ACCESS PATH IN FILE SYSTEM – READ()

6. Read the **lab4** inode and find pointers to its data blocks




# ACCESS PATH IN FILE SYSTEM – READ()

7. Read the **lab4** data blocks



# ACCESS PATH IN FILE SYSTEM – READ()

8. Update the **lab4** inode (e.g., last-accessed time)



# ACCESS PATH IN FILE SYSTEM

## – READ() TIMELINE

time ↓

	ROOT INODE	OS INODE	LAB4 INODE	ROOT DATA	OS DATA	LAB4 DATA
OPEN(LAB4)	READ	READ			READ	
READ()			READ		READ	
			READ	WRITE		READ

# OUTLINE

1. Introduction to Virtual File System (VFS)
2. Layout of osfs
3. Access Path in File System
4. Requirement & Grading

# REQUIREMENT

1. Complete **osfs\_create** function in dir.c (**7%**)
2. Complete **osfs\_write** function in file.c (**3%**)

# BONUS

1. Modify the osfs file allocation strategy (**2%**)  
e.g. (**extent-based** allocation or **multi-level indexing**)

# FUNCTIONS USE IN REQUIREMENT 1 -> CREATE()

## `osfs_create()`:

- `osfs_new_inode()`: Creates a new inode within the filesystem
- `osfs_add_dir_entry()`: Adds a new directory entry to a parent directory
- `d_instantiate()`: Associates a dentry with an inode in the VFS layer

# HINT FOR REQUIREMENT 1 -> CREATE()

- Step 1: Parse the parent directory passed by the VFS
- Step 2: Validate the file name length
- Step 3: Allocate and initialize a VFS & osfs inode
- Step 4: Parent directory entry update for the new file
- Step 5: Update the parent directory's metadata
- Step 6: Bind the inode to the VFS dentry

# FUNCTIONS USE IN REQUIREMENT 2 -> WRITE()

## `osfs_write()`:

- `file_inode()`: Retrieves the inode associated with an open file
- `osfs_alloc_data_block()`: Allocates a free data block from the block bitmap
- `copy_from_user()`: Copies data from user space to kernel space

# HINT FOR REQUIREMENT 2 -> WRITE()

- Step 1: Retrieve the inode and filesystem information
- Step 2: Check and allocate a data block if necessary
- Step 3: Limit the write length to fit within one data block
- Step 4: Write data from user space to the data block
- Step 5: Update inode & osfs\_inode attribute
- Step 6: Return the number of bytes written

# HINT FOR BONUS

**Data structure might be modified :**

`osfs_inode -> i_block`

`osfs_inode -> i_blocks`

`osfs_inode -> i_size`

**Functions might be modified:**

**All functions !!**

# HOW TO TEST & EXPECTED OUTPUT

1. After finishing the code, build the kernel module by running 'make'.
2. Use 'ls' to verify what files have been built

```
vboxuser@oslab:~/oslabfs$ ls  
dir.c Dockerfile file.c inode.c Makefile      Module.symvers osfs_init.c osfs.ko  osfs.mod.c osfs.o  super.o  
dir.o exec.sh    file.o inode.o modules.order  osfs.h        osfs_init.o osfs.mod  osfs.mod.o super.c
```

3. Insert the module into the Linux kernel: 'sudo insmod osfs.ko'
4. Run 'sudo dmesg' to check if the module is inserted correctly.

```
[ 637.738846] osfs: Successfully registered
```

5. Create a directory to serve as the mount point. In this example, the directory is named 'mnt'.

```
vboxuser@oslab:~/oslabfs$ ls  
dir.c Dockerfile file.c inode.c Makefile  modules.order  osfs.h       osfs_init.o  osfs.mod   osfs.mod.o  super.c  
dir.o exec.sh    file.o inode.o mnt      Module.symvers osfs_init.c  osfs.ko    osfs.mod.c  osfs.o    super.o
```

6. Mount the file system: 'sudo mount -t osfs none mnt/'
7. Run 'sudo dmesg' to check if the mount operation was successful.

```
[ 1026.834004] osfs: Filling super start  
[ 1026.834019] osfs_get_osfs_inode: Getting inode 1  
[ 1026.834021] osfs: Superblock filled successfully with root inode 1
```

# HOW TO TEST & EXPECTED OUTPUT

Now you are ready to test your write and create function

Enter the 'mnt' directory

1. Use 'sudo touch test1.txt' to test the create function. We expect that after running the command, a file named 'test1.txt' will appear in the 'mnt' directory.

```
vboxuser@oslab:~/oslabfs/mnt$ sudo touch test1.txt  
vboxuser@oslab:~/oslabfs/mnt$ ls  
test1.txt
```

2. Use 'sudo bash -c "echo 'I LOVE OSLAB' > test1.txt"' to test the write function. Then, run 'cat test1.txt'. The expected output is 'I LOVE OSLAB'.

```
vboxuser@oslab:~/oslabfs/mnt$ sudo bash -c "echo 'I LOVE OSLAB' > test1.txt"  
vboxuser@oslab:~/oslabfs/mnt$ ls  
test1.txt  
vboxuser@oslab:~/oslabfs/mnt$ cat test1.txt  
I LOVE OSLAB
```

